

Time-Frequency Analysis and Gabor Filtering of Both Simple and Complex Musical Samples

Trent Yarosevich¹

Abstract

We use windowed time-frequency analysis, specifically windowed Fourier transforms, to examine their utility and the impact of varying features such as window-size, filter-width and structure, as well as Low-Pass frequency filtering. The objective is to make high-level comments about the kinds of information these tools let us extract from music by examining it in a frequency domain.

https://github.com/tyarosevich/AMATH_582/blob/master/HW2/yarosevich_582_hw2_final.pdf

¹Department of Applied Mathematics, University of Washington, Seattle

Contents

1	Introduction and Overview	1
2	Theoretical Background	1
2.1	The Main Obstacle - Mathematical Uncertainty	1
2.2	The Solution - Windowed Transforms	2
2.3	Limitations	2
3	Algorithm Implementation and Development	2
3.1	Discretized Gabor Filtering	2
3.2	Sampling rate and Time Domain	2
3.3	Frequency Domain	3
	Fast Fourier Transform • Frequency Units	
3.4	Discretized Gabor Transform	3
	Step 1: slide vector and spectrogram matrix • Step 2: The window	
	• Step 3- Iteration • Step 4 - The Spectrogram • Step 5 - Overtone filtering	
4	Computational Results	4
4.1	Handel's "Messiah"	4
4.2	Part II - Mary Had a Little Lamb	4
5	Summary and Conclusions	5
	Appendix 1 - Matlab Functions	8
	Appendix 2 - Code Listing	8

1. Introduction and Overview

Music is comprised of oscillating signals formed by pressure waves in a medium (usually air), and thus it lends itself very naturally to analysis in a frequency domain. However, because we sense sound as the sum content of these pressure waves, the frequency content is not something we are always aware of. The interaction of these individual components is, however, a fundamental part of music (harmony).

Discussing the frequency content of music in the spatial domain requires expert knowledge of music theory. Changing to a frequency domain, however, let's us examine in detail

the frequency components of a given piece of music. The problem with this is that the evolution of those frequencies in time is lost. Time-frequency analysis is an approach that allows us to get some of the best of both worlds. Thus the basic high-level claim we want to support is that we can use time-frequency analysis to extract detailed frequency information from a piece of music while also knowing where it occurs in time. To illustrate this, the analysis will be comprised of two parts: in Part I, I will examine the famous Hallelujah Chorus from Handel's "Messiah", and use it to discuss the Gabor transform in general, as well as to study differently structured windows. In Part II, I will turn to a musically simple example, Mary Had A Little Lamb, in order to examine resolution in time and frequency, and explore the limitations of time-frequency analysis.

2. Theoretical Background

2.1 The Main Obstacle - Mathematical Uncertainty

The core of this analysis - and its obstacles - is the Fourier transform, which uses a kernel e^{-ikt} to change the basis functions of a given function $f(t)$ to sines and cosines. The transform and its inverse are as follows:

$$F(k) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-ikt} f(t) dt$$

$$f(t) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{ikt} F(k) dk$$
(1)

This change of bases uses integration to describe a spatial signal, in our case sound waves, in terms of their frequency components, but in doing so it loses information about momentum or position. Intuitively, this stems from the fact that as a change in time dt gets very small, the indeterminacy of a change in frequency k will get large - we can't detect a change in pitch in an infinitely small amount of time, which is why the Fourier transform integrates across the entire domain. Similarly, if determinacy of frequency is perfect, we would

be describing a perfect sine or cosine function with an infinite domain in time¹. From a purely mathematical point of view, we can observe that in the Fourier transform itself, all the dependence on time is integrated 'out' of the data.

2.2 The Solution - Windowed Transforms

The solution to this conundrum is to use 'windowed' transforms. Informally, this means that we can take the transform of a section of music localized in a certain span of time, thus giving a time-frame for its occurrence. One way of doing this is to use the Gábor Transform G , which is an intuitive extension of the Fourier transform that introduces a new kernel localizing information around a specific time τ ². Note that the function $g(\tau - t)$, which is used to localize a particular time, is assumed to be real and symmetric, with a 2-norm of 1:

$$g_{\tau,\omega}(\tau) = e^{i\omega\tau} g(\tau - t) \text{ the kernel}$$

$$G[f](t, \omega) = \bar{f}(t, \omega) = \int_{-\infty}^{\infty} f(\tau) \bar{g}(\tau - t) e^{-i\omega\tau} d\tau \quad (2)$$

One way to envision this is that the integration across all τ is 'sliding' the window across the signal, and extracting the frequency information inside each window. This then gives us the possibility of describing frequency content with respect to time, also known as a spectrogram. Given this resulting data, the original signal can then be recovered using the inverse transform. While we will not make use of this here, since the objective is to analyze data in the frequency domain, it is included for completeness as well as to make the instructive observation that the inverse transform is a double integral, unlike the inverse Fourier transform, since it must integrate across all frequency and time shifting components:

$$f(\tau) = \frac{1}{2\pi} \frac{1}{\|g\|^2} \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} \bar{f}_g(t, \omega) (\tau - t) e^{i\omega\tau} d\omega dt \quad (3)$$

2.3 Limitations

The fundamental problem of mathematical uncertainty is not avoided by the Gábor transform. If the 'window' that the transform slides across the signal is very narrow, we will recover frequency components at very specific timeframes, i.e. we will have good resolution in time. However, it should be very clear that this neglects lower frequency components whose period is significantly longer than the window. For example the window in Figure 1 would capture nearly all the frequencies present across about two seconds - very good resolution in frequency, but very poor in time (a lot happens in two seconds of music). Similarly, the opposite holds and a very wide window collects more frequency components, but is less localized in time. Illustrating this limitation and working around it will be one of the principle focuses of the analysis.




Figure 1. An illustration of a relatively 'wide' transform window

3. Algorithm Implementation and Development

In each section of this analysis, the same general algorithmic approach will be used several times, and so I will discuss its formulation once. In each case, data will be imported and processed to produce a spectrogram, with variations in parameters to conduct the analysis, but with the same overall structure each time.

3.1 Discretized Gábor Filtering

At the outset, we must note that because we are analyzing digital sound information, we must use a discretized Gábor transform, producing a lattice of time and frequency as shown in Figure 2³. This lattice occurs in uniform steps in both time




Figure 2. Motivation for discretized Fourier and Gábor Transforms

and frequency, which can be scale arbitrarily, but is indexed in m, n . The following discretized form of the transform is then used:

$$\bar{f}(m, n) = \int_{-\infty}^{\infty} f(t) \bar{g}_{m,n}(t) dt \quad (4)$$

3.2 Sampling rate and Time Domain

When the data is imported, it is an arbitrary vector of information, and in order to construct sensible domains in time and frequency, we have to know the sampling rate. In part 1, this rate is included with Matlab's built-in Handel file, but in Part II it is easily determined by dividing the length of the data vector by the duration, in seconds, of the sample. A time mesh can then be constructed by making an a vector with

¹"Feature Column from the AMS." American Mathematical Society, www.ams.org/publicoutreach/feature-column/fcarc-uncertainty.

²Kutz, Jose Nathan. Data-Driven Modeling and Scientific Computation. Methods for Complex Systems and Big Data. Oxford University Press, 2013.

³Ibid.

a number of points arbitrarily equal to the length of the data vector and dividing it by the sampling rate. This then gives suitably scaled time-mesh.

3.3 Frequency Domain

3.3.1 Fast Fourier Transform

In order to quickly process discretized data, we will be using the Fast Fourier Transform (FFT), which assumes a 2π periodic domain. Because this is a periodic domain, the last element of the spatial domain is omitted when constructing the K domain mesh, since this last element would be a repetition of the zero wave number.

3.3.2 Frequency Units

In Part I, I will simply scale the frequency domain mesh to a 2π periodic domain, so that the resulting domain is given in wave numbers k . This is accomplished by creating a mesh in time, and then noting the flipped domain of the FFT, which is then scaled by $\frac{2\pi}{L}$, as shown in the code-listing in **Appendix II** lines 26-7.

In Part II, however, we will be using a (Hz) based frequency domain since we will be considering musical pitch. In this case, we will then retain the scaling from the spatial domain and simply divide the frequency domain by the length L of the sample, to put the frequency in units of cycles per second.

3.4 Discretized Gábor Transform

The implementation of the discrete Gábor transform is relatively straightforward. In essence, we are using iteration to 'slide' a window across the data, as opposed to integrating in the continuous case.

3.4.1 Step 1: slide vector and spectrogram matrix

First, a vector is declared to determined how many windowed transforms are taken, and how far apart they are. Over and undersampling in this regard will be considered in the analysis, but in general, a larger number of smaller steps produces more accurate resolution in time. Second it is important to declare a matrix of zeros to hold the resulting data. The matrix rows correspond to the amplitude values across a frequency domain, and the columns to time. The number of rows is equal to the length of the tslide vector, and the columns to the length of the original time vector used to construct the frequency domain.

3.4.2 Step 2: The window

Next, the equation governing the window, $\bar{g}_{m,n}(t)$, must be chosen. There are many possible choices, and in the analysis we will use a Gaussian window, a so-called 'Mexican Hat' filter, and a Shannon filter. The construction of the Gaussian and Mexican Hat filters is fairly straightforward - their equations evaluated at a shifted time domain vector.

The Shannon filter is a little more complicated, but still fairly simple. To create a filter that steps up and then down after an arbitrary width, we essentially slide the t-mesh itself by the current value in the tslide vector, then apply this vector to a boolean evaluation of $t < \frac{w}{2}$ and $t > -\frac{w}{2}$. This returns a

boolean vector with zeros anywhere except a step centered around w which can then be multiplied by the data just as the other filters are. These filters are shown in Figure 3.




Figure 3. Plots of the filters used in the analysis

3.4.3 Step 3- Iteration

The actual iteration across the windows is done with a 'for' loop indexed from 1 to the length of the tslide vector. In each iteration we then evaluated the g function at the time mesh minus the current tslide value, multiply it component-wise with the data, and then take the FFT of the result and store it as a row of the holder matrix.

3.4.4 Step 4 - The Spectrogram

The resulting data is then used to create a spectrogram to visualize the frequency components over time. To make this more efficient, we can note that because we are transforming real spatial data, the negative frequencies will simply mirror the positive and can be discarded. In the case of Part II, we can take this a step further and observe that the frequency content of the samples falls in a small range. Thus by making observations from a plot in the frequency domain of the entire recording, we can ignore empty frequency domains by creating an index vector between two particular values. This is done by passing the shifted frequency mesh ks as an argument to the `find()` function in Matlab with boolean conditions for the upper and lower limit. The resulting vector can then be passed as an index when plotting the data. The data is then plotted using the `pcolor()` function in Matlab, which we use to plot the frequencies on the y-axis, time on the x-axis, and then color density using the matrix of amplitude data from the Gábor filtering.

3.4.5 Step 5 - Overtone filtering

The last algorithmic issue to discuss is the filtering of timbre. The timbre of an instrument is comprised of many elements, including overtones (integer multiples of the fundamental frequency of the note) as well as the structure of the fundamental frequency's peak, i.e. it's 'attack' and 'decay' as well as the small frequency components immediately around it. In the analysis of Part II, this timbre will be filtered out to examine the fundamental frequencies more closely.

To do this I simply identified the peak fundamental frequencies using a spatial plot of the sample, then centered three






Figure 4. The three low-pass filters for removing overtones



(a) Spatial Structure



(b) Frequency Structure
Figure 5

Gaussian filters around these peaks. These filters were then evaluated across the shifted frequency domain, added together as shown in Figure 4, and multiplied by the transformed and shifted data. The data was then shifted back to the FFT domain, inverse transformed, and Gábor filtered as described above. Note that the width of the Gaussian filters is controlled with parameter w , and a wide filter can be used to remove only the overtones, or a very narrow one to also filter out the attack and decay of the fundamental frequencies.

4. Computational Results

4.1 Handel's "Messiah"

To begin analyzing Handel's "Messiah" it is useful to motivate the need for frequency analysis by looking Figure 5 and simply observe what was noted at the outset: the spatial plot tells us little about the frequency content, and the frequency-domain plot tells us nothing about frequency components over time.

To begin a time-frequency analysis, we first consider two spectrograms with opposite goals - one uses a very narrow window in order to achieve excellent resolution in time, and

the other a wide window to achieve excellent resolution in space. In both cases, the window slides in increments of 0.1 seconds, and Figure 6 shows the dramatic difference in approach. In the frequency-resolved spectrogram we see the multitude of frequencies, clustered around 10 or more that are likely fundamental frequencies, and a large number of overtones. The time-resolved spectrogram on the other hand retains evidence of the cadence of the piece, with gaps between each "Hallelujah" evident in the graph.

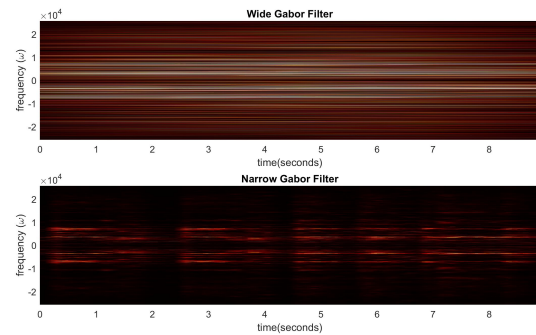


Figure 6. A narrow versus wide filter

The next objective was to compare the result of the different types of filters. The three are compared in Figure 7 with good resolution in time. The resulting spectrograms suggest that the Shannon filter is particularly good at achieving resolution in time while retaining more frequency information than the other filters, with an extreme example of a Shannon filter with a width of .1 - the same as the slide width - shown in Figure 8. In general, the different filters performed fairly similarly, however at moderately small window sizes, the Shannon filter again seems to give slightly better resolution in time. This can be seen with careful examination in figure 7, in which the the Shannon filter manages to capture the slight amplitude reduction between the first and second syllables of "Hallelujah".

I also investigated the effects of over and undersampling, which had the sensible result that large steps between the transform windows resulted in very poor resolution in time, giving the worst of both worlds - poor resolution in both time and frequency, as shown in Figure 9.

4.2 Part II - Mary Had a Little Lamb

This section narrowed in on simpler samples in order to analyze the frequency components themselves in a more qualitatively meaningful way than the "Messiah", which is harmonically complex. First, we considered the frequency domains of the two instruments, Piano and Recorder. Figure 10 shows clear differences. In particular, the piano has notable overtones, whereas the recorder has a slow attack, probably caused by the exhalation pressure of the player building to a peak. Figure 11 shows the various spectrograms produced with narrow and wide Gábor filters both before and after timbre filtering. As expected, the narrow windows give us good resolution in time, and we can see the individual notes

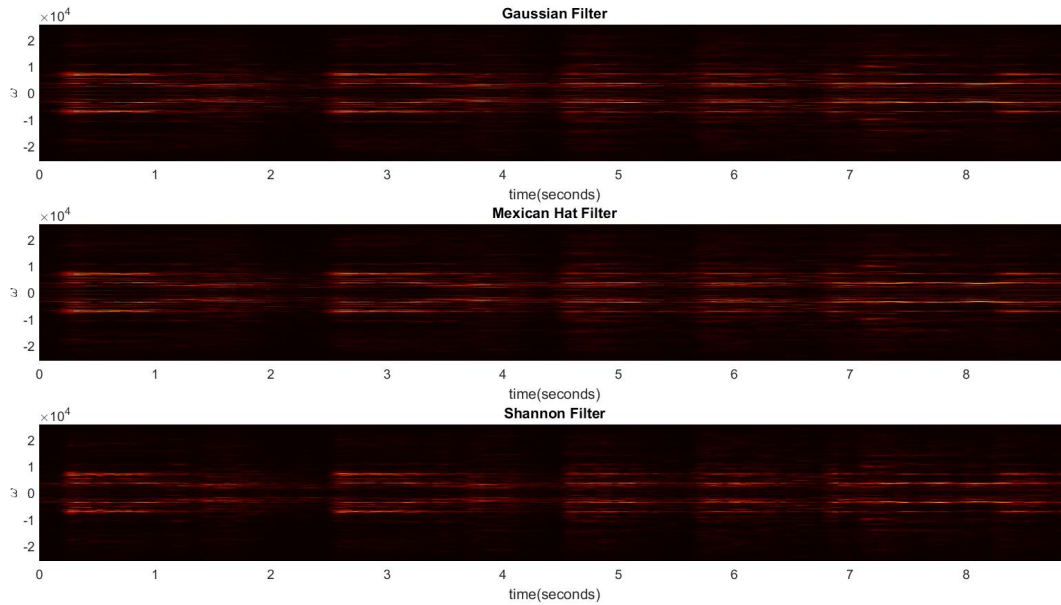

Figure 7. Comparison of different Gabor Filters g 

Figure 8. A finely time-resolved shannon filter

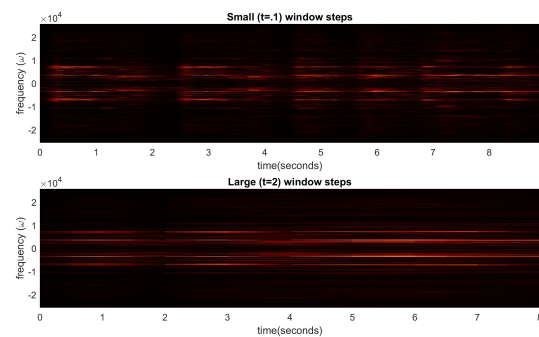


Figure 9. Under/over-sampling

of the melody. Even with a wide filter, however, it is a little difficult to see the fundamental frequencies due to the timbre. After timbre filtering, we then see exceptional resolution in frequency, particularly for the recorder.


5. Summary and Conclusions

The high-level goals of the analysis were satisfied, particularly in Part II. We saw that Gabor filtering was effective at isolating frequency information and charting its development over time. Perhaps the most important demonstration of this is found in Figure 6 in which we have attained resolution in time and frequency side-by-side, such that one could examine frequency content at a given time, and use the good time resolution spectrogram to associate it with specific portions of music, e.g. one could hone in on the frequency content of the "Ha" in "Hallelujah" if one were so inclined.


For example, even without the ubiquitous knowledge of

the melody and cadence of "Mary Had A Little Lamb", we can deduce from Figure 11(e) and (g) that the melody is playing notes approximately three times every two seconds, at frequencies of about 318 Hz, 287 Hz, and 256 (Hz) - namely, E4, D4, and C4 on a slightly out of tune (flat) Piano. Similarly, we can deduce that the Recorder is playing notes around 1035 Hz, 912 Hz, and around 820 Hz - which are somewhat close to C5 flat, B4 flat, and A4 flat.


In addition to the overall goals, I encountered a few issues that merit mention. For one, I found that using filters to take the original sample close to a pure tone was sometimes problematic, as removing too much of the signal could make any resolution in time difficult. This also wound up somewhat circular in its logic, since if we know what the fundamental frequency is, we could simply generate a pure tone anyway. Reducing a complex piece to pure tones might be helpful, but it would require some clever implementation to distinguish fundamentals from overtones.




(a) Piano in (Hz) domain.



(b) Recorder in (Hz) domain



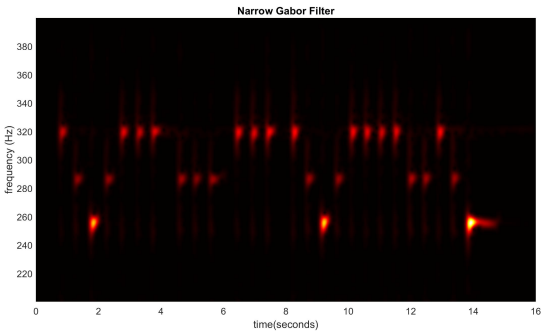
(c) Piano after Timbre Filtering




(d) Recorder after Timbre Filtering

Figure 10. Before and after Timbre Filtering

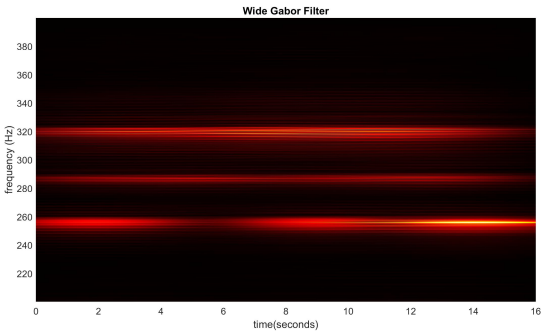
Another interesting development is found in the subtle timbre in the piano sample, which can be seen in 11(a). I am unsure what these frequencies represent, but they appear to occur around previously played notes and occur at precisely the same time as the currently played notes. I can only speculate that perhaps these are frequencies caused when the previously played strings in the piano are still vibrating slightly, and some sympathetic vibration from the currently played note kicks it up into barely detectable amplitude.



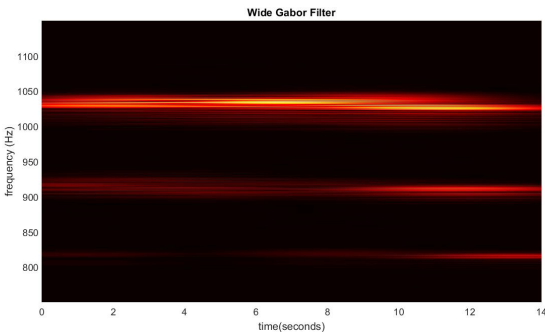
(a) Piano with narrow window



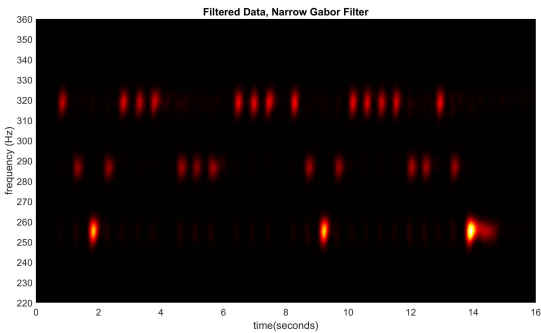
(b) Recorder with narrow window



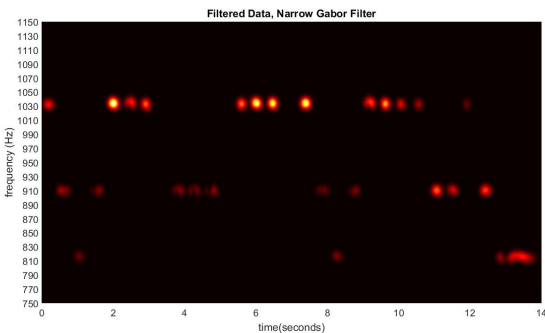
(c) Piano with wide window




(d) Recorder with wide window



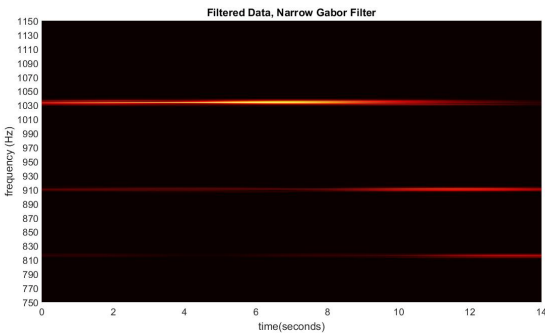
(e) Filtered piano, narrow window



(f) Filtered recorder, narrow window



(g) Filtered piano, wide window



(h) Filtered recorder, wide window

Figure 11. bla

Appendix 1

linspace(a,b,n): returns a vector in the span $[a,b]$ with n points.

fftshift(data)/ifftshift(data): Shifts data between zero-centered wave number domain and the domain used by the Fast Fourier Transform.

abs(data): Returns the absolute value of the passed object.

pcolor(a, b, c): Pcolor accepts two vectors which it plots on the x and y axes, and then accepts matrix of size a x b which it uses to convey a third dimension through color density in accordance with a variety of optional color maps.

Appendix 2

```

1  clc; clear all; close all;
2
3  % Load and plot the The Messiah segment
4
5  load handel
6   $v = y'/2;$ 
7  plot((1:length(v))/Fs,v);
8  xlabel("Time [sec]");
9  ylabel("Amplitude");
10 title("Handel's Messiah");
11 axis([0 9 0 1])
12 %% Plot with a 'window'
13  $t\_mesh = (1:\text{length}(v))/Fs;$ 
14 hold on
15 plot(t_mesh, exp(-(t_mesh - 4).^2), 'LineWidth',
16      2)
17 legend('Spatial Signal', "G\abor Window", '
18      Interpreter','latex')
19
20 %% Playback
21
22  $p8 = \text{audioplayer}(v,Fs);$ 
23 playblocking(p8);
24
25 %% Build the time and frequency domains
26
27  $L = \text{length}(v)/Fs;$   $n = \text{length}(v);$ 
28  $t2 = \text{linspace}(0, L, n+1);$   $t = t2(1:n);$ 
29  $k = (2*\pi/L)*[0:n/2-1 -n/2:-1];$   $ks = \text{fftshift}(k);$ 
30  $v = v(1:\text{end} - 1);$ 
31  $v\_t = \text{fft}(v);$ 
32
33 %% Plot in freq. domain of original file
34
35 close all
36 plot(ks, abs(fftshift(v_t))/max(abs(v_t)), 'k'); %
37  $\text{axis}([-50\ 50\ 0\ 1])$ 
38
39  $\text{set}(gca, 'FontSize', [14])$ 
40 xlabel('frequency (\omega)'), ylabel('Amplitude')
41
42 %% Gabor Window Iteration/Slide
43
44 close all
45  $\text{figure}(4)$ 
46  $tslide = 0:0.1:L;$ 
47
48  $\text{messiah\_spec}_1 = \text{zeros}(\text{length}(tslide), \text{length}(t)$ 
49  $-1);$ 

```



```

43 messiah_spec_2 = zeros(length(tslide), length(t)
    -1);
44
45 % Two different bands to compare
46 w1 = .1;
47 w2 = 100;
48
49 for j=1:length(tslide)
50     g=exp(-w1*(t(1:end)-tslide(j)).^2); % Gabor
51     v_g= g.*v;
52     v_g_t=fft(v_g);
53     messiah_spec_1(j, :) = abs(fftshift(v_g_t));
54 end
55 for j=1:length(tslide)
56     g=exp(-w2*(t(1:end)-tslide(j)).^2); % Gabor
57     v_g= g.*v;
58     v_g_t=fft(v_g);
59     messiah_spec_2(j,:) = abs(fftshift(v_g_t));
60 end
61
62 figure(5)
63 subplot(2,1,1)
64 pcolor(tslide,ks,messiah_spec_1. '),
65 shading interp
66 set(gca,'FontSize',[14])
67 colormap(hot)
68 xlabel('time(seconds)'), ylabel('frequency (\omega
    )')
69 title('Wide Gabor Filter');
70
71 subplot(2,1,2)
72 pcolor(tslide,ks,messiah_spec_2. '),
73 shading interp
74 set(gca,'FontSize',[14])
75 colormap(hot)
76 xlabel('time(seconds)'), ylabel('frequency (\omega
    )')
77 title('Narrow Gabor Filter')
78
79 %% Under versus oversampling
80 close all
81 %figure(4)
82 tslide1=0:0.1:L;
83 tslide2=0:1:L;
84 messiah_spec_1 = zeros(length(tslide1), length(t)
    -1);
85 messiah_spec_2 = zeros(length(tslide2), length(t)
    -1);
86
87
88 w = 100;
89
90 for j=1:length(tslide1)
91     g=exp(-w*(t(1:end)-tslide1(j)).^2); % Gabor
92     v_g= g.*v;
93     v_g_t=fft(v_g);
94     messiah_spec_1(j, :) = abs(fftshift(v_g_t));
95 end
96 for j=1:length(tslide2)
97     g=exp(-w*(t(1:end)-tslide2(j)).^2); % Gabor
98     v_g= g.*v;
99     v_g_t=fft(v_g);
100    messiah_spec_2(j,:) = abs(fftshift(v_g_t));
101 end
102
103 subplot(2,1,1)
104 pcolor(tslide1,ks,messiah_spec_1. '),
105 shading interp
106 set(gca,'FontSize',[14])
107 colormap(hot)

```

```

108 xlabel('time(seconds)'), ylabel('frequency (\omega
    )')
109 title('Small (t=.1) window steps');
110
111 subplot(2,1,2)
112 pcolor(tslide2,ks,messiah_spec_2. '),
113 shading interp
114 set(gca,'FontSize',[14])
115 colormap(hot)
116 xlabel('time(seconds)'), ylabel('frequency (\omega
    )')
117 title('Large (t=2) window steps')
118
119 %% Plot the different filters
120 close all
121
122 x_mesh = linspace(-3, 3, 100);
123 y_g = exp(-20*(x_mesh(1:end)).^2);
124 y_m = (1 - x_mesh.^2).*exp(-10*(x_mesh).^2);
125 t_shift = t(1:end-1)- 2;
126 y_sh = t_shift < 1/4 & t_shift > -1/4;
127
128 subplot(3,1,1)
129 plot(x_mesh, y_g)
130 axis([-3 3 -.2 1.2])
131 title("Gaussian Filter")
132
133 subplot(3,1,2)
134 plot(x_mesh, y_m)
135 axis([-3 3 -.1 1.2])
136 title("Mexican Hat Filter")
137
138
139 subplot(3,1,3)
140 plot(t(1:end-1), y_sh)
141 axis([0 4 -.2 1.2])
142 title("Shannon Filter")
143
144
145 %% Plot Spectrograms with Multiple types of
    filters
146 close all
147
148 tslide=0:0.05:L;
149 w = 100;
150
151 % Spatial resolution with the shannon filter is
    pretty outstanding
152 % with this filter width - the doubled 'hallelujah
    's are pretty clear
153 % in the spectrogram.
154 w_sh = .1;
155 messiah_sp_gauss = zeros(length(tslide), length(t)
    -1);
156 messiah_sp_mex = zeros(length(tslide), length(t)
    -1);
157 messiah_sp_shan = zeros(length(tslide), length(t)
    -1);
158
159 for j=1:length(tslide)
160     g=exp(-w*(t(1:end-1)-tslide(j)).^2);
161     v_g= g.*v;
162     v_g_t=fft(v_g);
163     messiah_sp_gauss(j,:)=abs(fftshift(v_g_t));
164 end
165
166 for j=1:length(tslide)
167     g=(1 - (t(1:end-1) - tslide(j)).^2).*exp(-w*(
        t(1:end-1)-tslide(j)).^2);
168     v_g= g.*v;

```

```

169     v_g_t=fft(v_g);
170     messiah_sp_mex(j,:)=abs(fftshift(v_g_t));
171 end
172
173 for j=1:length(tslide)
174
175     % Shift the t-mesh by the slide value, then
176     % use this
177     % to return a vector of booleans depending on
178     % whether or not
179     % the value in the t-mesh is greater than the
180     % -width/2 AND less
181     % than the width/2. This centers a square
182     % filter at tslide(j) or tau.
183     t_shift = t(1:end-1) - tslide(j);
184     g = t_shift < w_sh/2 & t_shift > -w_sh/2;
185     v_g = g.*v;
186     v_g_t=fft(v_g);
187     messiah_sp_shan(j,:)=abs(fftshift(v_g_t));
188 end
189
190 %% Plot the data using pcolor
191 close all
192
193 subplot(3,1,1)
194 pcolor(tslide,ks,messiah_sp_gauss. '),
195 shading interp
196 set(gca,'FontSize',[14])
197 colormap(hot)
198 xlabel('time(seconds)'), ylabel('\omega')
199 title('Gaussian Filter');
200
201 subplot(3, 1, 2)
202 pcolor(tslide,ks,messiah_sp_mex. '),
203 shading interp
204 set(gca,'FontSize',[14])
205 colormap(hot)
206 xlabel('time(seconds)'), ylabel('\omega')
207 title('Mexican Hat Filter')
208
209 subplot(3,1,3)
210 pcolor(tslide,ks,messiah_sp_shan. '),
211 shading interp
212 set(gca,'FontSize',[14])
213 colormap(hot)
214 xlabel('time(seconds)'), ylabel('\omega')
215 title('Shannon Filter')
216
217 %% Part II
218
219 clc; clear all; close all;
220
221 % tr_piano=16; % record time in seconds
222 % y_p=audioread('music1.wav'); Fs=length(y_p)/
223 % tr_piano;
224 % plot((1:length(y_p))/Fs,y_p);
225 % xlabel('Time [sec]'); ylabel('Amplitude');
226 % title('Mary had a little lamb (piano)'); drawnow
227 % p8 = audioplayer(y_p,Fs); playblocking(p8);
228 % figure(2)
229 tr_rec=14; % record time in seconds
230 y_p=audioread('music2.wav'); Fs=length(y_p)/tr_rec
231 ;
232 plot((1:length(y_p))/Fs,y_p);
233 xlabel('Time [sec]'); ylabel('Amplitude');
234 title('Mary had a little lamb (recorder)');
235 % p8 = audioplayer(y_p,Fs); playblocking(p8);
236
237 %%
238 L = 14; n = length(y_p);

```

```

233 t2 = linspace(0, L, n+1); t = t2(1:n);
234 % k = (2*pi/L)*[0:n/2-1 -n/2:-1]; ks = fftshift(k
    );
235 k = (1/L)*[0:(n/2-1) -n/2:-1];
236 ks = fftshift(k);
237
238 %% Plot the Piano in Frequency Domain
239 y_p_t = fft(y_p);
240 plot(ks, -abs(fftshift(y_p_t))/max(y_p_t), '
    LineWidth', 1)
241 axis([750 2200 -.1 1.1])
242 xticks([700 816 910 1034 1100 1200 1300 1400 1500
    1600 1700 1800 1900 2000 2100 2200])
243 xticklabels({'700', '816', '910', '1034', '1100',
    '1200', '1300', '1400', '1500', '1600', '1700',
    '1800', '1900', '2000', '2100', '2200'})
244 xlabel('Frequency (Hz)'), ylabel('Amplitude')
245 title('Recorder - Frequency Domain');
246 %% Check width of gabor filter
247 w = .1;
248 plot(t, exp(-w*(t(1:end)-8).^2))
249
250 %% Gabor Window Iteration/Slide
251 close all; clc;
252 %figure(4)
253 tslide=0:0.05:L;
254 piano_spec = zeros(length(tslide), length(t));
255 w = 100;
256 y_p = y_p';
257
258
259 for j=1:length(tslide)
260     g=exp(-w*(t(1:end)-tslide(j)).^2); % Gabor
261     y_g= g.*y_p;
262     y_g_t=fft(y_g);
263     piano_spec(j,:) = abs(fftshift(y_g_t));
264 end
265
266 % Range of frequency to plot, based on observation
    of the frequency domain
267 % plot.
268 idx = find(ks > 750 & ks < 1150);
269
270 pcolor(tslide, ks(idx), piano_spec(:, idx).')
271 shading interp
272 set(gca, 'FontSize', [14])
273 colormap(hot)
274 %axis([0 16 -.2 .2])
275 xlabel('time(seconds)'), ylabel('frequency (Hz)')
276 title('Narrow Gabor Filter');
277
278 %%
279 y_p_t_s = fftshift(y_p_t);
280 % [B, I] = maxk(y_p_t_s, 3);
281 % This wound up being annoying because the top
    three max freq were
282 % all around the first (loudest) note.
283
284 % Gaussian filters to remove all overtones. The k1
    /2/3 are
285 % simply found by examining a frequency domain
    plot of the sample.
286
287 % Piano Peaks
288 % k1 = 255;
289 % k2 = 287;
290 % k3 = 318;
291
292 % Recorder Peaks
293 k1 = 816;

```

```

294 k2 = 910;
295 k3 = 1034;
296
297 w = .1;
298
299 f1 = exp(-w*(ks - k1).^2);
300 f2 = exp(-w*(ks - k2).^2);
301 f3 = exp(-w*(ks - k3).^2);
302
303 y_p_filtered = (y_p_t_s.* (f1 + f2 + f3));
304 %% Plot the filters To confirm width
305 % Note the filter width matters, since a very
    narrow filter
306 % will drastically reduce the amplitude and
    produce less contrast
307 % in the pcolor plot with sound versus silence.
308 plot(ks, f1 + f2 + f3)
309 axis([750 1200 -1 1.1])
310 xlabel('time(seconds)'), ylabel('frequency (Hz)')
311 title('Low-Pass Filter');
312 %% Plot the filtered frequency domain
313
314 close all
315 plot(ks, -abs(y_p_filtered)/max(y_p_filtered))
316 axis([750 1100 -1 1.1])
317 xlabel('Frequency(Hz)'), ylabel('Amplitude')
318 title('Filtered Recorder - Freq. Domain');
319 %% Now a spectrogram of the filtered data
320 close all; clc;
321 %figure(4)
322 tslide=0:0.05:L;
323 filtered_spec = zeros(length(tslide), length(t));
324
325
326 % Two different bands to compare
327 w = .1;
328 w_sh = .1;
329
330 y_p_f = ifft(ifftshift(y_p_filtered));
331
332
333 for j=1:length(tslide)
334     g=exp(-w*(t(1:end)-tslide(j)).^2); % Gabor
335
336     % Shannon Filter (gaussian seems to work
        better)
337     % t_shift = t(1:end) - tslide(j);
338     % g = t_shift < w_sh/2 & t_shift > -w_sh/2;
339
340     y_g= g'.*y_p_f;
341     y_g_t=fft(y_g);
342     filtered_spec(j,:) = abs(fftshift(y_g_t));
343 end
344
345 %% Plot using pcolor
346
347 pcolor(tslide,ks(idx),filtered_spec(:,idx).')
348 shading interp
349 set(gca,'FontSize',[14], 'YLim', [750 1150], '
    YTick', [750:20:1150])
350 % set(gca,'FontSize',[14])
351 colormap(hot)
352 %axis([0 16 -2 .2])
353 xlabel('time(seconds)'), ylabel('frequency (Hz)')
354 title('Filtered Data, Narrow Gabor Filter');
355
356 % Results are perfect. The Piano is out of tune,
    ranging from around 5
357 % to 10 Hz flat.
358

```

```
359 %% Play the filtered data
360
361 % Note the playback sounds like pure tones, just
    as we would expect
362 p8 = audioplayer(y_p_f,Fs); playblocking(p8);
```